

Imitation-Projected Programmatic Reinforcement Learning



Abhinav Verma*
Rice University



Hoang M. Le*
Caltech



Yisong Yue
Caltech



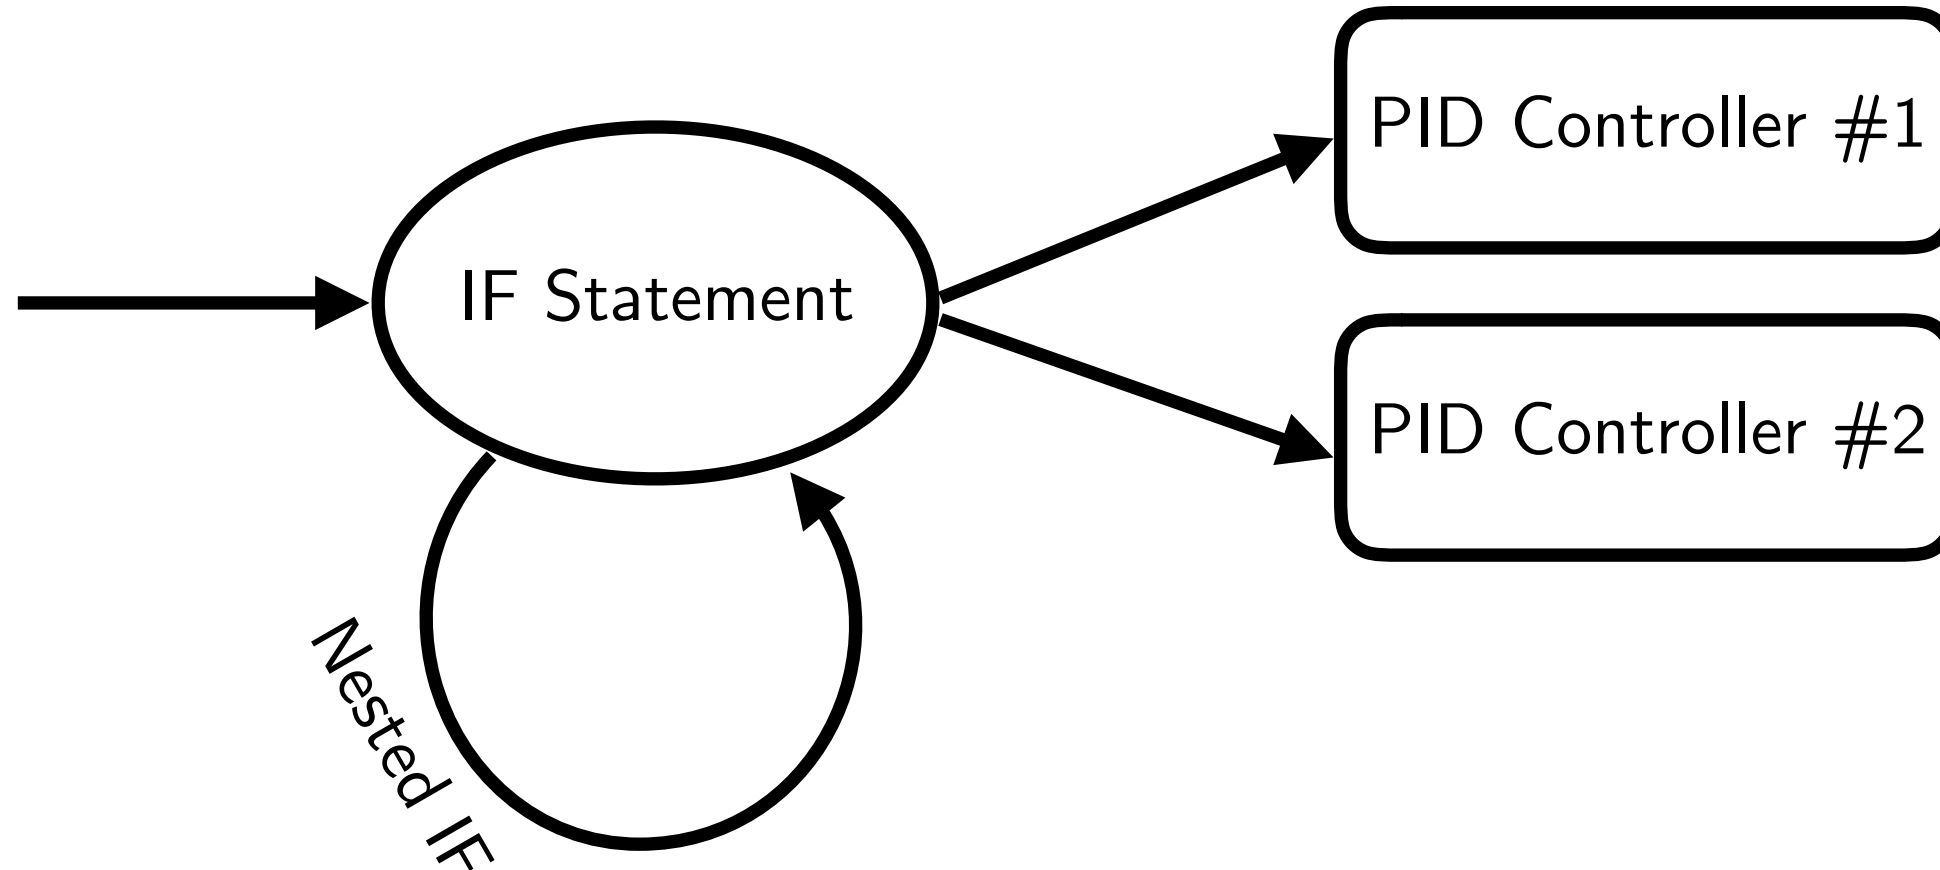
Swarat Chaudhuri
Rice University

(* equal contribution)

33rd Conference on Neural Information Processing Systems (NeuRIPS 2019)

Why learning programmatic policies

- Natural way to encode structure into policy class
- Benefits: policy-based guarantee, symbolic verification of programs (& interpretable)
- Example: domain-specific language over a family of parameterized simple controllers



Programmatic reinforcement learning

- The program space Π
 - language (arithmetic, boolean, relational) over simple policies

- **Goal:** find the best program

$$\pi^* = \operatorname{argmin}_{\pi \in \Pi} J(\pi)$$

- Learning programmatic policies is challenging:
 - Program Synthesis: highly structured nature of policy space
 - Deep RL: programmatic policy representation incompatible with deep policy search

- Approach Overview:

Building program structure into policy search via “lift-and-project”

Making use of hybrid representation

- Neural policy class F : deep RL representation
 - flexible, but unstable and does not satisfy desired property
- Programmatic policy class Π
 - less flexible, but certifiable
- Hybrid representation (functional regularization)

$$H \equiv \Pi \oplus F$$

$$h \equiv \pi + \lambda f \quad \text{defined as} \quad h(x) = \pi(x) + \lambda f(x)$$

Imitation-projected programmatic reinforcement learning (PROPEL)

hybrid class: $H \equiv \Pi \oplus F$


each iteration: $h_t \leftarrow \text{UPDATE}_F(\pi_{t-1})$

$\pi_t \leftarrow \text{PROJECT}_\Pi(h_t)$

UPDATE: $f \leftarrow f - \eta \lambda \nabla_F J(\pi + \lambda f)$

$h \leftarrow \pi + \lambda f$

PROJECT: imitation learning



Analysis: approximate (functional) mirror descent perspective

hybrid class: $H \equiv \Pi \oplus F$

each iteration: $h_t \leftarrow \text{UPDATE}_F(\pi_{t-1}) \approx \text{UPDATE}_H(\pi_{t-1})$


$\pi_t \leftarrow \text{PROJECT}_\Pi(h_t) \approx \text{argmin}_{\pi \in \Pi} \|\pi - h_t\|^2$

UPDATE: $f \leftarrow f - \eta \lambda \nabla_F J(\pi + \lambda f)$

$h \leftarrow \pi + \lambda f$

$\text{UPDATE}_H(\pi_{t-1}) = \pi_{t-1} - \nabla_H J(\pi_{t-1})$

PROJECT_Π : finite sample error



Experiment in TORCS

Program space: decision tree with simple, parameterized PID controllers at leaf nodes



“if the car is aligned with the axis of the track...”

if ($\text{obs}_{\text{TrackPos}}(0) < 0.001$ and $\text{obs}_{\text{TrackPos}}(0) > -0.001$)
 then $PID_{\text{rpm}}(0.44, 4.92, 0.89, 49.79)$
 else $PID_{\text{rpm}}(0.40, 4.92, 0.89, 49.79)$

“then accelerate, otherwise slow down”

Experiment in TORCS - Results

Performance: Lap time in seconds / Crash-ratio (over 25 seeds, lower is better)

LENGTH	G-TRACK 3186M	E-ROAD 3260M	AALBORG 2588M	RUUDSKOGEN 3274M	ALPINE-2 3774M
PRIOR	312.92 / 0.0	322.59 / 0.0	244.19 / 0.0	340.29 / 0.0	402.89 / 0.0
DDPG	78.82 / 0.24	89.71 / 0.28	101.06 / 0.4	CR / 0.68	CR / 0.92
NDPS	108.25 / 0.24	126.80 / 0.28	163.25 / 0.4	CR / 0.68	CR / 0.92
VIPER	83.60 / 0.24	87.53 / 0.28	110.57 / 0.4	CR / 0.68	CR / 0.92
PROPELPROG	93.67 / 0.4	119.17 / 0.4	147.28 / 0.12	124.58 / 0.16	256.59 / 0.16
PROPELTREE	78.33 / 0.4	79.39 / 0.4	109.83 / 0.16	118.80 / 0.24	236.01 / 0.36

Generalization: PROPEL completed 10/20 unseen tracks, DDPG completed 2/20

	G-TRACK	E-ROAD	AALBORG	RUUDSKOGEN	ALPINE-2
G-TRACK	-	124 / CR	CR / CR	CR / CR	CR / CR
E-ROAD	102 / 92	-	CR / CR	CR / CR	CR / CR
AALBORG	201 / 91	228 / CR	-	217 / CR	CR / CR
RUUDSKOGEN	131 / CR	135 / CR	CR / CR	-	CR / CR
ALPINE-2	222 / CR	231 / CR	184 / CR	CR / CR	-

Code available at: <https://bitbucket.org/averma8053/propel>